

Bayesian Learning of Sum-Product Networks

Martin Trapp^{1,2} R. Peharz³ H. Ge³
F. Pernkopf¹ Z. Ghahramani^{4,3}

¹Graz University of Technology, ²OFAI

³University of Cambridge, ⁴Uber AI

October 29, 2019

- ▶ Sum-product networks (SPNs) [Poon2011] are flexible density estimators and have received significant attention due to their attractive inference properties.
- ▶ While parameter learning in SPNs is well developed, most structure learners are somewhat adhoc and based on intuition rather than a clear learning principle.

- ▶ Sum-product networks (SPNs) [Poon2011] are flexible density estimators and have received significant attention due to their attractive inference properties.
- ▶ While parameter learning in SPNs is well developed, most structure learners are somewhat adhoc and based on intuition rather than a clear learning principle.
- ▶ We introduce a well-principled Bayesian framework for SPN structure learning by decomposing the problem into:
 1. laying out a computational graph, and
 2. learning the so-called scope function over the graph.

- ▶ Sum-product networks (SPNs) [Poon2011] are flexible density estimators and have received significant attention due to their attractive inference properties.
- ▶ While parameter learning in SPNs is well developed, most structure learners are somewhat adhoc and based on intuition rather than a clear learning principle.
- ▶ We introduce a well-principled Bayesian framework for SPN structure learning by decomposing the problem into:
 1. laying out a computational graph, and
 2. learning the so-called scope function over the graph.
- ▶ We propose a natural parametrisation for an important and widely used special case of SPNs, incorporate the parameters into a Bayesian model, and perform posterior inference over structure and parameters jointly.

Sum-Product Networks

What is a Sum-Product Network?

- ▶ Let $\mathbf{X} = \{X_1, \dots, X_D\}$ be set of D random variables, for which N i.i.d. samples are available.
- ▶ An SPN is a distribution over \mathbf{X} defined as a 4-tuple $\mathcal{S} = (\mathcal{G}, \psi, w, \theta)$.

What is a Sum-Product Network?

- ▶ Let $\mathbf{X} = \{X_1, \dots, X_D\}$ be set of D random variables, for which N i.i.d. samples are available.
- ▶ An SPN is a distribution over \mathbf{X} defined as a 4-tuple $\mathcal{S} = (\mathcal{G}, \psi, w, \theta)$.
 - \mathcal{G} is a computational graph.

What is a Sum-Product Network?

- ▶ Let $\mathbf{X} = \{X_1, \dots, X_D\}$ be set of D random variables, for which N i.i.d. samples are available.
- ▶ An SPN is a distribution over \mathbf{X} defined as a 4-tuple $\mathcal{S} = (\mathcal{G}, \psi, w, \theta)$.
 - \mathcal{G} is a computational graph.
 - ψ is a so-called scope function.

What is a Sum-Product Network?

- ▶ Let $\mathbf{X} = \{X_1, \dots, X_D\}$ be set of D random variables, for which N i.i.d. samples are available.
- ▶ An SPN is a distribution over \mathbf{X} defined as a 4-tuple $\mathcal{S} = (\mathcal{G}, \psi, w, \theta)$.
 - \mathcal{G} is a computational graph.
 - ψ is a so-called scope function.
 - w denotes the set of sum-weights and θ the set of leaf node parameters.

What is a Sum-Product Network?

- ▶ Let $\mathbf{X} = \{X_1, \dots, X_D\}$ be set of D random variables, for which N i.i.d. samples are available.
- ▶ An SPN is a distribution over \mathbf{X} defined as a 4-tuple $\mathcal{S} = (\mathcal{G}, \psi, w, \theta)$.
 - \mathcal{G} is a computational graph.
 - ψ is a so-called scope function.
 - w denotes the set of sum-weights and θ the set of leaf node parameters.
- ▶ Note: This definition is conceptually different to the classic definition of SPNs as it disentangles the definition of the SPNs structure into a computational graph, which has only few requirements, and a scope function, which ensures completeness and decomposability.

Computational Graph \mathcal{G}

- ▶ Is a connected directed acyclic graph (DAG), containing three types of nodes: sums (S), products (P) and leaves (L).
- ▶ Only encodes the topological layout of the nodes, while the effective SPN structure is encoded via the scope function ψ .

Computational Graph \mathcal{G}

- ▶ Is a connected directed acyclic graph (DAG), containing three types of nodes: sums (S), products (P) and leaves (L).
- ▶ Only encodes the topological layout of the nodes, while the effective SPN structure is encoded via the scope function ψ .




Figure: Example of a tree-shaped computational graph with two layers.

Scope Function ψ

- ▶ Let \mathbf{N} denote the set of all nodes.
- ▶ ψ a function $\psi: \mathbf{N} \mapsto 2^{\mathbf{X}}$ assigning each node in the graph a sub-set of \mathbf{X} ($2^{\mathbf{X}}$ denotes the power set of \mathbf{X}).

Scope Function ψ

- ▶ Let \mathbf{N} denote the set of all nodes.
- ▶ ψ a function $\psi: \mathbf{N} \mapsto 2^{\mathbf{X}}$ assigning each node in the graph a sub-set of \mathbf{X} ($2^{\mathbf{X}}$ denotes the power set of \mathbf{X}).

It has the following properties:

1. If N is the root node, then $\psi(N) = \mathbf{X}$.
2. If N is a sum or product, then $\psi(N) = \bigcup_{N' \in \text{ch}(N)} \psi(N')$.
3. For each $S \in \mathbf{S}$ we have
 $\forall N, N' \in \text{ch}(S): \psi(N) = \psi(N')$ (*completeness*).
4. For each $P \in \mathbf{P}$ we have
 $\forall N, N' \in \text{ch}(P): \psi(N) \cap \psi(N') = \emptyset$ (*decomposability*).

Scope Function ψ


- ▶ Let \mathbf{N} denote the set of all nodes.
- ▶ ψ a function $\psi: \mathbf{N} \mapsto 2^{\mathbf{X}}$ assigning each node in the graph a sub-set of \mathbf{X} ($2^{\mathbf{X}}$ denotes the power set of \mathbf{X}).

It has the following properties:


1. If N is the root node, then $\psi(N) = \mathbf{X}$.
2. If N is a sum or product, then $\psi(N) = \bigcup_{N' \in \text{ch}(N)} \psi(N')$.
3. For each $S \in \mathbf{S}$ we have
 $\forall N, N' \in \text{ch}(S): \psi(N) = \psi(N')$ (*completeness*).
4. For each $P \in \mathbf{P}$ we have
 $\forall N, N' \in \text{ch}(P): \psi(N) \cap \psi(N') = \emptyset$ (*decomposability*).

Note: Completeness and decomposability are necessary for any SPN to be a well-defined probability distribution and to allow exact inference in linear time (in the model size).

Example SPN $\mathcal{S} = (\mathcal{G}, \psi, w, \theta)$



↓ Apply scope function ψ on \mathcal{G}



Bayesian Learning of Sum-Product Networks

Bayesian Parameter Learning

- ▶ The key insight for Bayesian parameter learning [Zhao2016, Rashwan2016, Vergari2019] is that sum nodes can be interpreted as latent variables Z_S , clustering data instances.

Bayesian Parameter Learning

- The key insight for Bayesian parameter learning [Zhao2016, Rashwan2016, Vergari2019] is that sum nodes can be interpreted as latent variables Z_S , clustering data instances.

$$\begin{aligned}
 \mathcal{S}(\mathbf{x}) &= \sum_{\mathbf{z}} \prod_{S \in \mathbf{S}} w_{S, \mathbf{z}_S} \prod_{L \in T(\mathbf{z})} L(\mathbf{x}_L) \\
 &= \sum_{\mathcal{T}} \prod_{(S, N) \in \mathcal{T}} w_{S, N} \prod_{L \in \mathcal{T}} L(\mathbf{x}_L) \underbrace{\left(\sum_{\bar{\mathbf{z}}} \prod_{S \in \bar{\mathbf{S}}_{\mathcal{T}}} w_{S, \bar{\mathbf{z}}_S} \right)}_{=1} \quad (1)
 \end{aligned}$$

Bayesian Parameter Learning

- The key insight for Bayesian parameter learning [Zhao2016, Rashwan2016, Vergari2019] is that sum nodes can be interpreted as latent variables Z_S , clustering data instances.

$$\begin{aligned}
 \mathcal{S}(\mathbf{x}) &= \sum_{\mathbf{z}} \prod_{S \in \mathcal{S}} w_{S, \mathbf{z}_S} \prod_{L \in \mathcal{T}(\mathbf{z})} L(\mathbf{x}_L) \\
 &= \sum_{\mathcal{T}} \prod_{(S, N) \in \mathcal{T}} w_{S, N} \prod_{L \in \mathcal{T}} L(\mathbf{x}_L) \underbrace{\left(\sum_{\bar{\mathbf{z}}} \prod_{S \in \bar{\mathcal{S}}_{\mathcal{T}}} w_{S, \bar{\mathbf{z}}_S} \right)}_{=1} \quad (1)
 \end{aligned}$$

- \mathcal{T} is a so-called induced tree [Zhao2016] which is a sub-tree in \mathcal{S} such that the root of \mathcal{S} is the root of \mathcal{T} , each $S \in \mathcal{T}$ has only one child and each $P \in \mathcal{T}$ has the same children as in \mathcal{S} .
- $\mathcal{T}(\mathbf{z})$ is a surjective (not injective) function that assigns to each value \mathbf{z} the induced tree \mathcal{T} determined by \mathbf{z} .

Bayesian Parameter Learning (cont.)

- It is now conceptually straightforward to extend an SPN to a Bayesian setting, by equipping the sum-weights and leaf-parameters with suitable priors.

Bayesian Parameter Learning (cont.)

- It is now conceptually straightforward to extend an SPN to a Bayesian setting, by equipping the sum-weights and leaf-parameters with suitable priors.

Generative model for Bayesian parameter learning:

$$\begin{aligned} w_S \mid \alpha &\sim \mathcal{Dir}(w_S \mid \alpha) \quad \forall S, & z_{S,n} \mid w_S &\sim \mathcal{Cat}(z_{S,n} \mid w_S) \quad \forall S \forall n, \\ \theta_L \mid \gamma &\sim p(\theta_L \mid \gamma) \quad \forall L, & \mathbf{x}_n \mid \mathbf{z}_n, \theta &\sim \prod_{L \in T(\mathbf{z}_n)} L(\mathbf{x}_{L,n} \mid \theta_L) \quad \forall n. \end{aligned} \tag{2}$$

Bayesian Structure & Parameter Learning

- ▶ We introduce Bayesian learning of SPNs by considering the case in which \mathcal{G} is a tree-shaped region graph.

Bayesian Structure & Parameter Learning

- ▶ We introduce Bayesian learning of SPNs by considering the case in which \mathcal{G} is a tree-shaped region graph.
- ▶ A region graph \mathcal{R} can be understood as a vectorised representation of SPNs and is a connected DAG containing two types of nodes: regions ($R \in \mathbf{R}$) and partitions ($P \in \mathbf{P}$).




Figure: Example region-graph. Based on the illustration by [Peharz2019].

Bayesian Structure & Parameter Learning (cont.)

- For each data dimension d we introduce a discrete latent variable $Y_{P,d}$.

Bayesian Structure & Parameter Learning (cont.)

- ▶ For each data dimension d we introduce a discrete latent variable $Y_{P,d}$.
- ▶ The latent variable represents a decision to assign dimension d to a particular child, given that all partitions above decided to assign d onto the (unique) path leading to the partition

Bayesian Structure & Parameter Learning (cont.)

- ▶ For each data dimension d we introduce a discrete latent variable $Y_{P,d}$.
- ▶ The latent variable represents a decision to assign dimension d to a particular child, given that all partitions above decided to assign d onto the (unique) path leading to the partition

Generative model for joint Bayesian learning:

$$\begin{aligned}
 w_S \mid \alpha &\sim \text{Dir}(w_S \mid \alpha) \quad \forall S, & z_{S,n} \mid w_S &\sim \text{Cat}(z_{S,n} \mid w_S) \quad \forall S \forall n, \\
 v_P \mid \beta &\sim \text{Dir}(v_P \mid \beta) \quad \forall P, & y_{P,d} \mid v_P &\sim \text{Cat}(y_{P,d} \mid v_P) \quad \forall P \forall d, \\
 \theta_L \mid \gamma &\sim p(\theta_L \mid \gamma) \quad \forall L, & \mathbf{x}_n \mid \mathbf{z}_n, \mathbf{y}, \theta &\sim \prod_{L \in T(\mathbf{z}_n)} L(\mathbf{x}_{\mathbf{y},n} \mid \theta_L) \quad \forall n.
 \end{aligned}
 \tag{3}$$

$\mathbf{x}_{\mathbf{y},n}$ denotes the evaluation of L on the scope induced by \mathbf{y} .

Posterior Inference

- We perform Gibbs sampling alternating between i) updating parameters w , θ (fixed y), and ii) updating y (fixed w , θ) to learn Bayesian SPNs.

Posterior Inference

- ▶ We perform Gibbs sampling alternating between i) updating parameters w, θ (fixed y), and ii) updating y (fixed w, θ) to learn Bayesian SPNs.
- ▶ This approach has shown to be sufficient for most real-world dataset, but more sophisticated approaches, e.g. particle Gibbs combined with Hamiltonian Monte Carlo sampling; variational inference or posterior bootstrap, might be a interesting future avenues for large-scale problems.

Experiments (discrete data)

Dataset	LearnSPN	RAT-SPN	CCCP	ID-SPN	ours	ours [∞]	BTD
NLTCS	-6.11	-6.01	-6.03	-6.02	-6.00	-6.02	-5.97
MSNBC	-6.11	-6.04	-6.05	-6.04	-6.06	-6.03	-6.03
KDD	-2.18	-2.13	-2.13	-2.13	-2.12	-2.13	-2.11
Plants	-12.98	-13.44	-12.87	-12.54	-12.68	-12.94	-11.84
Audio	-40.50	-39.96	-40.02	-39.79	-39.77	-39.79	-39.39
Jester	-53.48	-52.97	-52.88	-52.86	-52.42	-52.86	-51.29
Netflix	-57.33	-56.85	-56.78	-56.36	-56.31	-56.80	-55.71
Accidents	-30.04	-35.49	-27.70	-26.98	-34.10	-33.89	-26.98
Retail	-11.04	-10.91	-10.92	-10.85	-10.83	-10.83	-10.72
Pumsb-star	-24.78	-32.53	-24.23	-22.41	-31.34	-31.96	-22.41
DNA	-82.52	-97.23	-84.92	-81.21	-92.95	-92.84	-81.07
Kosarak	-10.99	-10.89	-10.88	-10.60	-10.74	-10.77	-10.52
MSWeb	-10.25	-10.12	-9.97	-9.73	-9.88	-9.89	-9.62
Book	-35.89	-34.68	-35.01	-34.14	-34.13	-34.34	-34.14
EachMovie	-52.49	-53.63	-52.56	-51.51	-51.66	-50.94	-50.34
WebKB	-158.20	-157.53	-157.49	-151.84	-156.02	-157.33	-149.20
Reuters-52	-85.07	-87.37	-84.63	-83.35	-84.31	-84.44	-81.87
20 Newsgrp	-155.93	-152.06	-153.21	-151.47	-151.99	-151.95	-151.02
BBC	-250.69	-252.14	-248.60	-248.93	-249.70	-254.69	-229.21
AD	-19.73	-48.47	-27.20	-19.05	-63.80	-63.80	-14.00

Experiments (heterogeneous data)

Dataset	MSPN	ABDA	ours	ours [∞]
Abalone	9.73	2.22	3.92	3.99
Adult	-44.07	-5.91	-4.62	-4.68
Australian	-36.14	-16.44	-21.51	-21.99
Autism	-39.20	-27.93	-0.47	-1.16
Breast	-28.01	-25.48	-25.02	-25.76
Chess	-13.01	-12.30	-11.54	-11.76
Crx	-36.26	-12.82	-19.38	-19.62
Dermatology	-27.71	-24.98	-23.95	-24.33
Diabetes	-31.22	-17.48	-21.21	-21.06
German	-26.05	-25.83	-26.76	-26.63
Student	-30.18	-28.73	-29.51	-29.9
Wine	-0.13	-10.12	-8.62	-8.65

Experiments (missing data)




Figure: EachMovie (D: 500, N: 5526) **Figure:** WebKB (D: 839, N: 3361)




Figure: BBC (D: 1058, N: 1895)

Conclusion

- ▶ Structure learning is an important topic in SPNs, and many promising directions have been proposed in recent years.
- ▶ However, most of these approaches are based on intuition and refrain from declaring an explicit and global principle to structure learning.
- ▶ In this paper, our primary motivation is to *change* this practice

Conclusion

- ▶ Structure learning is an important topic in SPNs, and many promising directions have been proposed in recent years.
- ▶ However, most of these approaches are based on intuition and refrain from declaring an explicit and global principle to structure learning.
- ▶ In this paper, our primary motivation is to *change* this practice
- ▶ We phrase structure (and joint parameter) learning as Bayesian inference in a latent variable model.
- ▶ In various experiments we show that this principled approach competes well with prior art and that we gain several benefits, such as *automatic protection against overfitting*, *robustness under missing data* and a natural extension to nonparametric formulations.

Conclusion

- ▶ Structure learning is an important topic in SPNs, and many promising directions have been proposed in recent years.
- ▶ However, most of these approaches are based on intuition and refrain from declaring an explicit and global principle to structure learning.
- ▶ In this paper, our primary motivation is to *change* this practice
- ▶ We phrase structure (and joint parameter) learning as Bayesian inference in a latent variable model.
- ▶ In various experiments we show that this principled approach competes well with prior art and that we gain several benefits, such as *automatic protection against overfitting*, *robustness under missing data* and a natural extension to nonparametric formulations.
- ▶ A critical insight for our approach is to decompose structure learning into: constructing a computational graph and separately learning the SPN's scope function.

References

- [Poon2011] H. Poon and P. Domingos. Sum-product networks: A new deep architecture. In Proceedings of UAI, pages 337–346, 2011.
- [Zhao2016] H. Zhao, T. Adel, G. J. Gordon, and B. Amos. Collapsed variational inference for sum-product networks. In Proceedings of ICML, pages 1310–1318, 2016.
- [Rashwan2016] A. Rashwan, H. Zhao, and P. Poupart. Online and distributed Bayesian moment matching for parameter learning in sum-product networks. In Proceedings of AISTATS, pages 1469–1477, 2016.
- [Vergari2019] A. Vergari, A. Molina, R. Peharz, Z. Ghahramani, K. Kersting, and I. Valera. Automatic Bayesian density analysis. In Proceedings of AAAI, 2019.
- [Peharz2019] R. Peharz, A. Vergari, K. Stelzner, A. Molina, X. Shao, M. Trapp, K. Kersting, and Z. Ghahramani. Random sum-product networks: A simple but effective approach to probabilistic deep learning. In Proceedings of UAI, 2019.
- [Trapp2019] M. Trapp, R. Peharz, H. Ge, F. Pernkopf, and Z. Ghahramani. Bayesian Learning of Sum-Product Networks. To appear at NeurIPS, 2019.